
Zero-Ping: A Packet-Loss-Resilient Neural Codec for Real-Time Voice Communication

May 23, 2026

Luca Brunetti 1967993

Abstract

Who has not experienced the frustration of packet loss while playing online games? This project addresses packet loss concealment for real-time mono 16 kHz speech transmission in a "neural way". Zero-Ping is an end-to-end trained neural audio codec based on an RVQ-GAN architecture, designed to reconstruct the audio waveform even when some transmitted packets are lost.

1. Introduction

Unreliable transmission conditions are common in online gaming voice chat and, more generally, in real-time communication scenarios. Existing neural audio codecs are mainly designed to optimize compression efficiency and perceptual reconstruction quality under reliable transmission. However, in real-time communication, lost packets cannot be recovered by retransmission without introducing additional delay. As a result, packet loss may permanently remove part of the transmitted representation, severely compromising speech intelligibility.

In this project, I designed and trained Zero-Ping, a neural speech codec based on an RVQ + Transformer architecture, following the RVQ-GAN paradigm used by modern neural audio codecs. The natural transmission cycle is that the encoder receives a clean waveform, tokenizes it, and sends the resulting packets. During transmission, some packets can be lost. As a consequence, the decoder either receives no information for the missing time interval or simply continues decoding from the next packet that arrives. Realistical packet loss patterns do not follow a purely random distribution, such as an iid Bernoulli process. Instead, they typically occur in bursts over time. To simulate this behavior, I implemented the **Gilbert-Elliott**

model (1). This model describes packet loss through a two-state Markov process, where the network alternates between a good state and a bad state. The persistence of the bad state makes consecutive packet losses more likely, thus reproducing bursty loss behavior:




Figure 1. *Top-Figure*: Spectrogram of the original speech signal. *Bottom-Figure*: Packet-loss pattern generated with the Gilbert-Elliott model, with an overall loss rate of 17.9%. Green vertical lines mark lost packets, where each packet corresponds to a 15 ms audio frame.

Crucially, no pre-corrupted dataset is needed: the *Gilbert-Elliott* simulator is applied online, directly in the latent domain after the RVQ quantizer, at every training step. This mirrors the actual transmission scenario: the encoder quantizes and "sends" the latent frames, the channel drops some of them, and the decoder must reconstruct from what arrives. This ensures that the model is exposed to a virtually unlimited variety of loss scenarios during training, naturally covering the full range of burst lengths and loss rates that the Gilbert-Elliott needed can produce, without any offline data augmentation pipeline. Zero-Ping handles these time holes by predicting the missing latent vectors before they reach the decoder. Thus, it is capable of reconstructing speech, even in critical connection scenarios, without compromising intelligibility.

2. Related works

Although existing neural audio codecs are not primarily designed for packet loss concealment, **EnCodec** (2) and **SoundStream** (3) served as the main architectural refer-

Email: Luca Brunetti <brunetti.1967993@studenti.uniroma1.it>.

ences for Zero-Ping. Despite targeting different bitrate ranges and sampling frequencies, both systems inspired the real-time encoder–quantizer–decoder structure.

Additional implementation details were inspired by **Improved RVQGAN** (4), including RVQ codebook initialization, the use of Snake activations to improve waveform reconstruction, and several training design aspects.

3. Architecture

The project’s target is voice chat, so we operate at 16kHz. I decided to operate at a low bitrate (6kbps). Some priors that naturally driven the design: we know where and when a packet is missing, transmission protocols (like **RTP**(5)) includes sequence numbers and consistent timestamps. The overall system must satisfy the real-time constraint of one-way delay below 150 ms (**ITU-T G.114**(6) standard). Components:

- **Encoder:** A fully causal convolutional network mapping a mono waveform $\mathbf{x} \in \mathbb{R}^{1 \times T}$ to a latent sequence $\mathbf{z} \in \mathbb{R}^{128 \times T'}$, where $T' = T/240$. It consists of four downsampling stages, each preceded by dilated residual blocks to widen the receptive field, followed by a single-layer LSTM at the bottleneck and a linear projection to 128 dimensions. The resulting frame rate is $16000/240 \approx 66.7$ fps, so each latent frame corresponds to exactly 15 ms of audio (i assume one frame per transmitted packet).
- **Quantizer:** The latent $\mathbf{z} \in \mathbb{R}^{128}$ is quantized by a cascade of $N_Q = 9$ codebooks, each of size 1024. At each stage q , the nearest entry is found via cosine similarity and subtracted from the current residual (4). The quantized output is $\mathbf{z}_q = \sum_{q=1}^9 \mathbf{e}_q$. Since each codebook requires $\log_2 1024 = 10$ bits per frame and the frame rate is ≈ 66.7 fps, the resulting bitrate is $66.7 \times 10 \times 9 \approx 6$ kbps.
- **Repair Transformer:** A small Transformer that performs packet-loss concealment in the RVQ latent domain. Missing frames are replaced by a learned placeholder embedding; a binary mask embedding signals received vs. lost to the model. We exploit the positional prior given by RTP sequence numbers: we know exactly *which* frames are missing, so lost keys are hard-masked out of attention. The model attends over a causal buffer of 8 past frames and 2 future frames, adding 30 ms of lookahead latency. Crucially, once a lost frame is estimated it is committed to the buffer, so past context contains only decoded or previously repaired frames.
- **Decoder:** The inverse of the encoder but with *snake* activations, optimal for reconstructing waveforms.

The critical part on designing the architecture was the Real-Time/reconstruction quality tradeoff. Each component was tested with the **RTF-benchmark** before trained in order to stay inside the 150ms constraint counting up the 30ms lookahead and another 30ms of trasmission delay. The model was trained for circa 1M steps, first Enc-codebooks-Dec only, then Trasformer only and a final joint tune. Different datasets are used to capture various tones and microphone conditions (LibriTTS, VCTK, ecc...). For the training objectives at each "stage" check the Appendix-5.

4. Evaluation

Real-Time Performance: Benchmarks on an NVIDIA RTX 3050 Laptop GPU show that Zero-Ping exhibits a nearly fixed per-call overhead of approximately 70–75 ms, largely independent of the processed chunk duration. As a result, short chunks are dominated by invocation overhead rather than by the intrinsic computational cost of the model. At 75 ms, Zero-Ping reaches real-time performance on average, with $\text{RTF} = 0.996$, making this the minimum chunk size for average real-time execution on this hardware.

Table 1. (6 kbps). $\text{RTF} < 1.0$ indicates real-time capability.

Chunk	RTF	$\pm \text{std}$	Mean (ms)	Real-Time
60 ms	1.3164	0.2465	79.0	✗
75 ms	0.9960	0.1964	74.7	✓
90 ms	0.8173	0.1513	73.6	✓

Reconstruction Quality: The most common benchmark for quality is the **MUSHRA**, but since for me it was unpractical, i opted for objective metrics: STOI and PESQ. Figure 2 shows an example on a single audio signal.




Figure 2. Spectral error maps showing the effect of packet loss on the same speech audio shown in Figure 1. Both examples use the same Gilbert–Elliott packet-loss pattern. *Top*: reconstruction without the repair module, where packet losses introduce strong localized spectral distortions and a lower intelligibility score, $\text{STOI} = 0.841$. *Bottom*: Zero-Ping reconstruction with latent repair, which substantially reduces packet-loss artifacts and improves intelligibility, reaching $\text{STOI} = 0.942$. Color intensity represents the spectral difference in dB between the reconstructed and reference signal.

Final evaluation is done comparing ZPCodec against three baselines at the same 6kbps bitrate: **Opus** (7), the standard real-time VoIP codec; **EnCodec** (2), designed for low-latency streaming; and **DAC** (4), a high-quality neural codec that employs non-causal convolutions and is therefore not suitable for real-time deployment. It is included as a quality upper bound for neural coding at this bitrate. All codecs are subjected to the exact same Gilbert-Elliott packet-loss patterns.

Table 2. Benchmark on Mozilla Common Voice test-set of 3994 samples; dataset not seen in train. **Loss** is the % of lost packets.

Loss	Codec	PESQ	P-Win	STOI	S-Win
10%	ZPCodec	2.0920	50.33%	0.9017	60.48%
	DAC	2.0874	48.90%	0.8913	39.42%
	EnCodec	1.7558	0.73%	0.8510	0.10%
	Opus	1.5780	0.05%	0.8174	0.00%
20%	ZPCodec	1.8463	81.77%	0.8756	91.06%
	DAC	1.5882	17.45%	0.8166	8.79%
	EnCodec	1.4617	0.78%	0.7804	0.15%
	Opus	1.3744	0.00%	0.7520	0.00%

5. Conclusions

In the end, the evaluation shows that Zero-Ping consistently outperforms all baselines across both perceptual quality and speech intelligibility, with the gap widening as loss rate increases. Apart from audio chat, the proposed approach extends naturally to any low-bitrate audio transmission scenario where packet loss is frequent, such as live streaming over congested networks or audio delivery in weak-signal environments, where waveform-level concealment consistently falls short. The full project is available at <https://github.com/Lucabr01/Zero-Ping>.

Statement on the use of AI

The core usage was the exploration of literature, with strong focus on understanding all the parts related to signal/audio. Overall help with the code where i failed or needed support on implementing things like: the "connection" logic on quantization $\rightarrow$ transformer with the masks and the two-pass strategy (to fill the past embedding buffer at each "step"); tuning "mechanical" stuff for GPU optimization during training, for example in the joint training i used a RTX 5090 that required some code modifications. So in the architecture side, i entirely designed it and used AI just to help implementing whenever i needed support.

AI was heavily used on the evaluation side. In particular, it supported the implementation of the RTF benchmark

and helped in understanding and interpreting the obtained results. During the benchmark implementation on the *Mozilla Common Voice* test set, AI was also heavily used to support audio preprocessing, format conversion, batching, and the adaptation of the evaluation pipeline to the different codec APIs (Opus, EnCodec, and DAC). In the notebook that can be used to try the model, the slider was entirely made by AI.

In the report i used it only for the mathematical formalism on the appendix.

References

- [1] G. Haßlinger and O. Hohlfeld, "The gilbert-elliott model for packet loss in real time services on the internet," in *Proceedings of the 14th GI/ITG Conference on Measurement, Modelling and Evaluation of Computer and Communication Systems*, 2008, pp. 269–283.
- [2] A. Défossez, J. Copet, G. Synnaeve, and Y. Adi, "High fidelity neural audio compression," *arXiv preprint arXiv:2210.13438*, 2022. [Online]. Available: <https://arxiv.org/abs/2210.13438>
- [3] N. Zeghidour, A. Luebs, A. Omran, J. Skoglund, and M. Tagliasacchi, "Soundstream: An end-to-end neural audio codec," *arXiv preprint arXiv:2107.03312*, 2021. [Online]. Available: <https://arxiv.org/abs/2107.03312>
- [4] R. Kumar, P. Seetharaman, A. Luebs, I. Kumar, and K. Kumar, "High-fidelity audio compression with improved rvqgan," *arXiv preprint arXiv:2306.06546*, 2023. [Online]. Available: <https://arxiv.org/abs/2306.06546>
- [5] H. Schulzrinne, S. Casner, R. Frederick, and V. Jacobson, "RTP: A transport protocol for real-time applications," RFC 3550, Jul. 2003. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc3550>
- [6] ITU-T, "One-way transmission time," International Telecommunication Union, Recommendation G.114, May 2003. [Online]. Available: <https://www.itu.int/rec/T-REC-G.114>
- [7] J.-M. Valin, K. Vos, and T. Terriberry, "Definition of the opus audio codec," Internet Engineering Task Force, RFC 6716, Sep. 2012. [Online]. Available: <https://www.rfc-editor.org/rfc/rfc6716>

Appendix

Loss Functions

Time-domain loss.

$$\mathcal{L}_{\text{time}} = \frac{1}{T} \sum_{t=1}^T |x_t - \hat{x}_t| \quad (1)$$

L1 loss on the raw waveform. Directly penalises amplitude errors sample by sample.

RVQ commitment loss.

$$\mathcal{L}_{\text{commit}} = \frac{1}{N_Q} \sum_{q=1}^{N_Q} \left(1 - \frac{\mathbf{r}_q \cdot \text{sg}(\mathbf{e}_q)}{\|\mathbf{r}_q\| \|\text{sg}(\mathbf{e}_q)\|} \right) \quad (2)$$

where $\text{sg}(\cdot)$ denotes stop-gradient, $\mathbf{e}_q$ is the nearest codebook entry at residual step q , and $\mathbf{r}_q$ is the current residual. Encourages the encoder to produce embeddings aligned with codebook entries in direction rather than magnitude, consistent with the cosine-similarity lookup used at quantisation time.

Latent repair loss.

$$\mathcal{L}_{\text{repair}} = \frac{1}{|\mathcal{M}| \cdot D} \sum_{t \in \mathcal{M}} \|\mathbf{z}_t^{\text{pre}} - \mathbf{z}_t^{\text{post}}\|_1 \quad (3)$$

where $\mathcal{M} = \{t : m_t = 0\}$ is the set of lost frames, $\mathbf{z}^{\text{pre}}$ is the original RVQ latent (target) and $\mathbf{z}^{\text{post}}$ is the Repair Transformer estimate. The mask restricts supervision to lost frames only; on received frames the two tensors coincide by construction.

Discriminator hinge loss.

$$\mathcal{L}_D = \frac{1}{K} \sum_{k=1}^K \left(\mathbb{E}[\text{ReLU}(1 - D_k(x))] + \mathbb{E}[\text{ReLU}(1 + D_k(\hat{x}))] \right) \quad (4)$$

where $K = 4$ is the number of STFT scales (windows $\in \{128, 256, 512, 1024\}$ samples) and D_k is the patch discriminator at scale k . Encourages real logits ≥ 1 and fake logits ≤ -1 .

Multi-scale mel loss.

$$\mathcal{L}_{\text{mel}} = \frac{1}{S} \sum_{s=1}^S \left(\|\mathbf{M}_s(x) - \mathbf{M}_s(\hat{x})\|_1 + \log(\mathbf{M}_s(x) + \epsilon) - \log(\mathbf{M}_s(\hat{x}) + \epsilon) \right) \quad (5)$$

where $S = 7$ STFT scales (window lengths $\in \{32, \dots, 2048\}$) and $\mathbf{M}_s$ is the amplitude mel-spectrogram at scale s . The linear term emphasises energetic peaks (formants, plosives); the log term preserves weaker components (fricatives).

Generator adversarial loss.

$$\mathcal{L}_{\text{adv}} = -\frac{1}{K} \sum_{k=1}^K \mathbb{E}[D_k(\hat{x})] \quad (6)$$

Pushes the generator to produce reconstructions that the discriminator scores as real.

Feature matching loss.

$$\mathcal{L}_{\text{fm}} = \frac{1}{KL} \sum_{k=1}^K \sum_{l=1}^L \|\phi_k^{(l)}(x) - \phi_k^{(l)}(\hat{x})\|_1 \quad (7)$$

where $\phi_k^{(l)}$ denotes the l -th intermediate feature map of discriminator k . Real feature maps are detached, so gradients flow only through the generator. Stabilises GAN training by encouraging the generator to match the internal representations of real audio at every layer.

Training Stages

Encoder-Quantizer-Decoder only: The encoder, RVQ, and decoder are trained jointly with the full generator objective $\mathcal{L}_G = \mathcal{L}_{\text{time}} + \lambda_{\text{mel}} \mathcal{L}_{\text{mel}} + \lambda_{\text{commit}} \mathcal{L}_{\text{commit}} + \lambda_{\text{adv}} \mathcal{L}_{\text{adv}} + \lambda_{\text{fm}} \mathcal{L}_{\text{fm}}$. The Repair Transformer is not involved. This stage establishes a stable latent space and trains the discriminator before any concealment is introduced.

Repair Transformer only: The codec and discriminator weights are frozen. Only the Repair Transformer is trained using $\lambda_{\text{rep}} \mathcal{L}_{\text{repair}} + \lambda_{\text{mel}} \mathcal{L}_{\text{mel}}$, where $\mathcal{L}_{\text{repair}}$ supervises the reconstructed latents on lost frames only and $\mathcal{L}_{\text{mel}}$ provides an auditory quality signal on the decoded output. Freezing the codec ensures the latent space is not perturbed while the repair module learns to exploit it.

Joint fine-tuning: All components are unfrozen and trained end-to-end with the same objective as Stage 1, i.e. $\mathcal{L}_G$ without $\mathcal{L}_{\text{repair}}$. Dropping the latent supervision encourages the system to optimise perceptual quality directly, allowing the encoder, decoder, and Repair Transformer to adapt jointly under the adversarial objective.